

# Analysis of a Priority Queue for Task Scheduling

Abhishek Vishwakarma

June 2025

## **Abstract**

This report details the design choices, implementation, and performance analysis of a Priority Queue data structure built for a simulated task scheduling application. The implementation utilizes a binary min-heap represented by a Python list and supports dynamic task insertion, extraction of highest- and lowest-priority tasks, and priority updates. A thorough analysis of the time complexity for each core operation is provided, justifying the efficiency of the chosen approach.

## Contents

<b>1</b>	<b>Design Choices and Implementation</b>	<b>3</b>
1.1	Data Structure for the Binary Heap . . . . .	3
1.2	The Task Class . . . . .	3
1.3	Choice of a Min-Heap . . . . .	3
<b>2</b>	<b>Analysis of Core Operations</b>	<b>3</b>
2.1	insert(task) . . . . .	3
2.2	extract_min() . . . . .	3
2.3	extract_max() . . . . .	4
2.4	decrease_key(task, new_priority) . . . . .	4
2.5	is_empty() . . . . .	4
<b>3</b>	<b>Analysis of Scheduling Results</b>	<b>4</b>

# 1 Design Choices and Implementation

## 1.1 Data Structure for the Binary Heap

A standard Python `list` was chosen to represent the binary heap. This choice is primarily justified by its direct compatibility with Python’s built-in `heapq` library, which provides highly optimized heap operations that work directly on lists. This approach offers several advantages:

- **Ease of Implementation:** It abstracts away the manual calculations for parent and child node indices, simplifying the code for heap manipulation.
- **Efficiency:** The underlying C implementation of the `heapq` module is very fast. Using a list leverages the efficiency of contiguous memory blocks, which is ideal for the implicit tree structure of a heap.

## 1.2 The Task Class

A dedicated `Task` class was designed to encapsulate all information related to a single task. It stores a `task_id`, a `description`, and a `priority`. The class implements the “less-than” special method (`__lt__`) to allow Python’s `heapq` to compare tasks based on their priority value, which is essential for maintaining the heap property.

## 1.3 Choice of a Min-Heap

The implementation uses a **min-heap**, where the element with the smallest value is always at the root. In the context of our scheduler, this means a **lower priority number corresponds to a higher actual priority**. For example, a task with priority ‘1’ will be processed before a task with priority ‘2’. This is a common and intuitive convention for task scheduling systems and aligns perfectly with the default behavior of the `heapq` library.

# 2 Analysis of Core Operations

The efficiency of the priority queue is determined by the time complexity of its core operations.

## 2.1 `insert(task)`

This operation adds a new task to the queue. The `heapq.heappush` function first adds the element to the end of the list and then “sifts it up” the heap, swapping it with its parent until the min-heap property is restored.

- **Time Complexity:**  $O(\log n)$ . In a balanced binary heap of  $n$  elements, the height is  $\log n$ . The sifting process involves at most  $\log n$  swaps.

## 2.2 `extract_min()`

This operation removes and returns the task(s) with the highest priority (lowest priority number). It uses `heapq.heappop`, which swaps the root element with the last element, removes the last element, and then “sifts down” the new root to restore the heap property. The logic is extended to pop all tasks that share the same minimum priority level.

- **Time Complexity:**  $O(k \log n)$ , where  $k$  is the number of tasks with the same highest priority. The initial ‘heappop’ is  $O(\log n)$ , and each subsequent pop for items with the same priority is also  $O(\log n)$ . For extracting a single item, this is simply  $O(\log n)$ .

### 2.3 `extract_max()`

Since the data structure is a min-heap, it is not optimized for finding the maximum element. The implemented `extract_max` function must iterate through all active entries in the heap to find the task with the highest priority number (lowest priority).

- **Time Complexity:**  $O(n)$ . This linear scan is required because any element in the heap could potentially be the maximum. After finding the max-priority tasks, they are marked for removal.

### 2.4 `decrease_key(task, new_priority)`

Updating an element's key in a heap can be complex. The chosen implementation uses a common and practical technique:

1. The old entry in the heap is marked as invalid (a "lazy removal").
2. A new entry with the updated priority is inserted into the heap.

This avoids a costly search-and-sift operation within the heap structure.

- **Time Complexity:**  $O(\log n)$ . The cost is dominated by the insertion of the new task.

### 2.5 `is_empty()`

This operation checks if there are any active tasks in the queue, ignoring any entries that have been marked as removed.

- **Time Complexity:**  $O(n)$ . In the worst case, this requires a full scan of the heap if all entries have been marked as removed but not yet popped.

## 3 Analysis of Scheduling Results

The behavior of the priority queue was tested through an interactive session, which confirmed its correctness and highlighted key design trade-offs.

- **Queue Behavior with Different Priorities:** Tasks were consistently managed according to their priority. When the 'Print all tasks' option was used, the list was correctly sorted, showing tasks with priority '1' first, followed by '2', and then '3', validating the core min-heap property.
- **Handling of Ties by `extract_min()`:** The scheduler correctly handled cases with multiple tasks at the highest priority. For instance, when tasks **Coding** (ID 2) and **Studying** (ID 3) both had priority 1, the `extract_min()` operation successfully identified and removed both of them in a single call. This demonstrates that the implementation correctly processes all tasks at a given priority level before moving to the next.
- **Handling of Ties by `extract_max()`:** Similar to the above, the scheduler proved robust in handling ties for the lowest-priority tasks. When **Cycling** (ID 1) and **Gaming** (ID 5) both had the lowest priority of 3, the `extract_max()` operation correctly found and removed both tasks simultaneously.
- **Implementation Trade-offs:** The interactive session revealed the practical consequences of the chosen implementation strategy. The "lazy removal" approach for updating and deleting keys is simple and keeps the `insert` and `decrease_key` operations efficient at  $O(\log n)$ . However, this simplicity comes at a cost:

- The `extract_max()` and `is_empty()` operations both have a time complexity of  $O(n)$ . This is because they must perform a linear scan over all entries to find the maximum value or to check for any remaining active tasks among the "removed" placeholders.
- For an application that frequently needs to find the lowest-priority task, this  $O(n)$  operation would be a significant bottleneck. A more complex but efficient solution would involve maintaining a second, parallel max-heap, though this would double the memory usage and complicate the insertion and deletion logic.

The current implementation is therefore optimized for applications where insertions and high-priority extractions are the most common operations, which is a typical use case for a task scheduler.